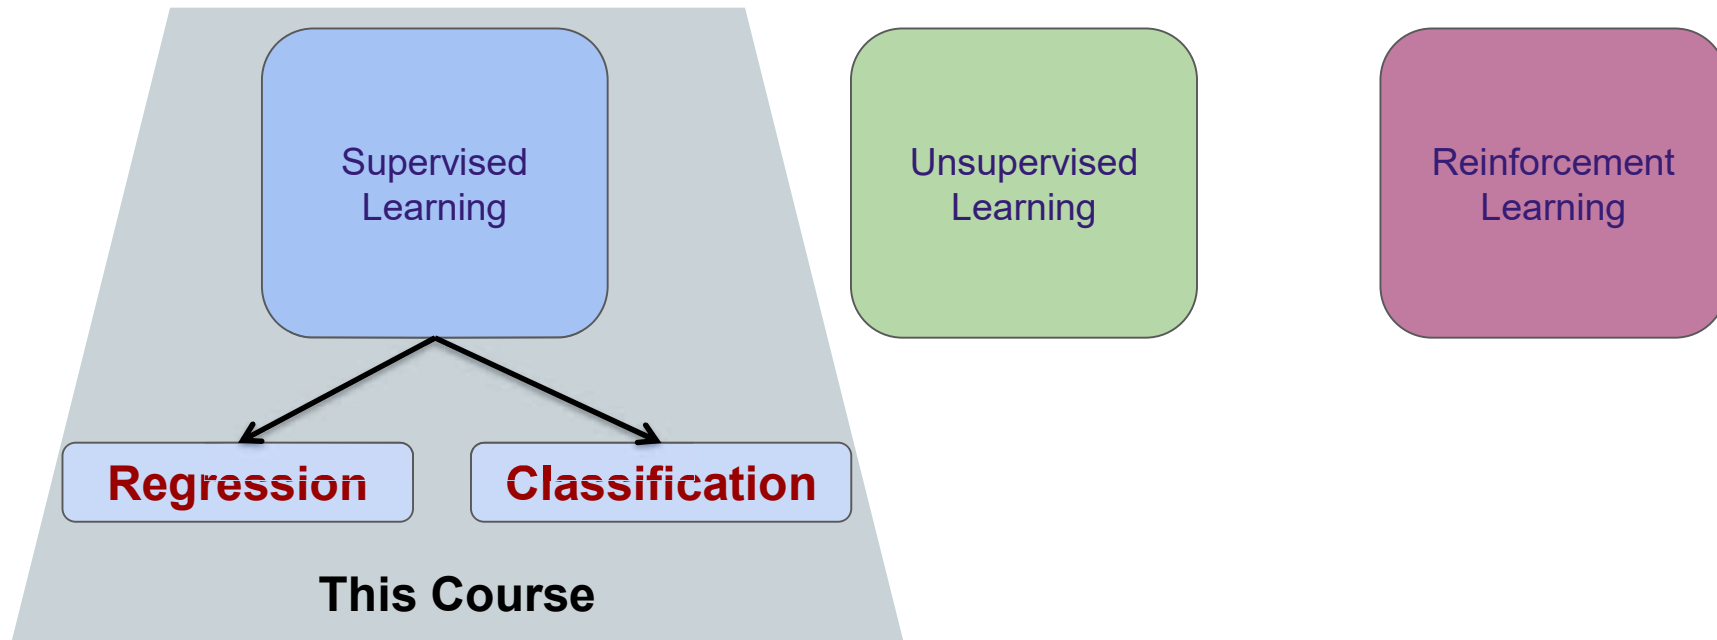


Course overview

1. Introduction to Classification
2. Linear Regression for Classification ?
3. Logistic Regression Model (Binary Classification)
 1. Logistic Regression Model
 2. Logistic Regression Cost function
4. Multi-Class Logistic Regression
5. Evaluating Classifiers
6. Cross Validation

Reminder of Machine Learning Types

- Machine learning tasks are typically classified into **three broad categories**

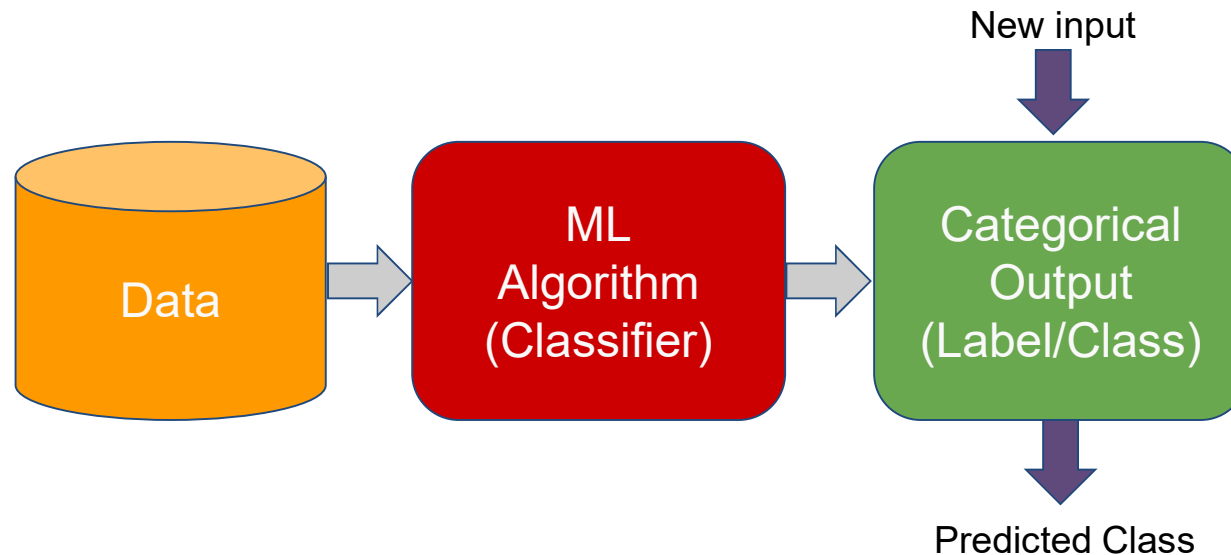


3.1

Introduction to Classification

Classification

- **Goal:** Inputs are divided into two or more classes, and the ML algorithm must produce a model that assigns unseen inputs to one or more of these classes
- An algorithm that implements classification is known as a **classifier**




Two-class(Binary) Classification

- **Emails type:** Output y has 2 categories



Input: x

Mail sender, subject, keywords, etc...



Output: y

Email type

Two-class(Binary) Classification

- **Loan demand:** Output y has 2 categories



Input: x

Client's characteristics
(age, Revenue, credit, etc..)



Safe

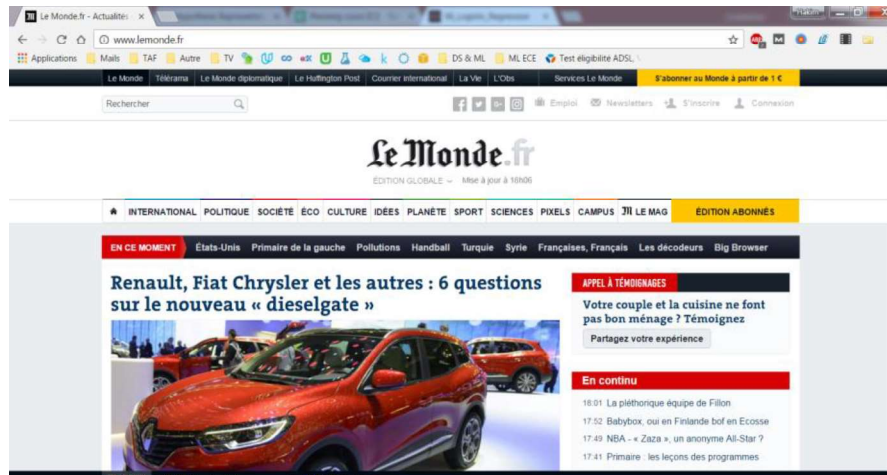
Risky

Output: y

Loan safety evaluation


Multi-class Classifier

- **News** : Output y has more than 2 categories



Input: x

Webpage: title, keywords, etc..




Output: y

News category

Multi-class Classifier

- **Weather:** Output y has more than 2 categories



Classification Algorithms

Linear Classifiers

Logistic Regression

Naive Bayes classifier

Linear discriminant

Support vector machines

Decision Trees

Random Forests

Boosting

Quadratic Classifiers

Neural Networks

K-nearest neighbor


Classification problem example

Can Linear Regression help ?

- Classification problem: Loan demand safety ?



Input: x
Client's characteristics
(age, Revenue, charges, etc..)



Output: y
Loan safety evaluation

Can Linear Regression help ?

- The classification problem is just like the regression problem
- Except that the **y values** to be predicted are discrete values.
 - Example: Loan demand evaluation: Safe or risky ?

Age	Revenue	Charges	...	Loan safety
20	1200	500	...	0
23	1700	450	...	0
25	1900	400	...	0
27	2000	450	...	0
40	2500	250	...	1
42	2750	200	...	1
45	2750	300	...	1
47	3000	100	...	1

x_1

x_2

x_3

y

$y \in \{0, 1\}$

→ $y = 1$: Positive Class

→ $y = 0$: Negative Class



Age	Loan safety
20	0
23	0
25	0
27	0
40	1
42	1
45	1
47	1

x

y

Safe

Risky


3.2

Linear Regression for classification ?

Can Linear Regression help ?

- One method is to use linear regression

x	y
Age	Loan safety
20	0
23	0
25	0
27	0
40	1
42	1
45	1
47	1



- Map all predictions greater than 0.5 as a 1 and all less than 0.5 as a 0
- Threshold classifier output $h\theta(x)$ at 0.5 :
 - if $h\theta(x) \geq 0,5 \rightarrow \text{predict: } y="1"$
 - if $h\theta(x) < 0,5 \rightarrow \text{predict: } y="0"$

Can Linear Regression help ?


- Linear Regression can sometimes be luckybut it is often not useful for classification problems

1st

- Another problem: is that classification need categorical values:

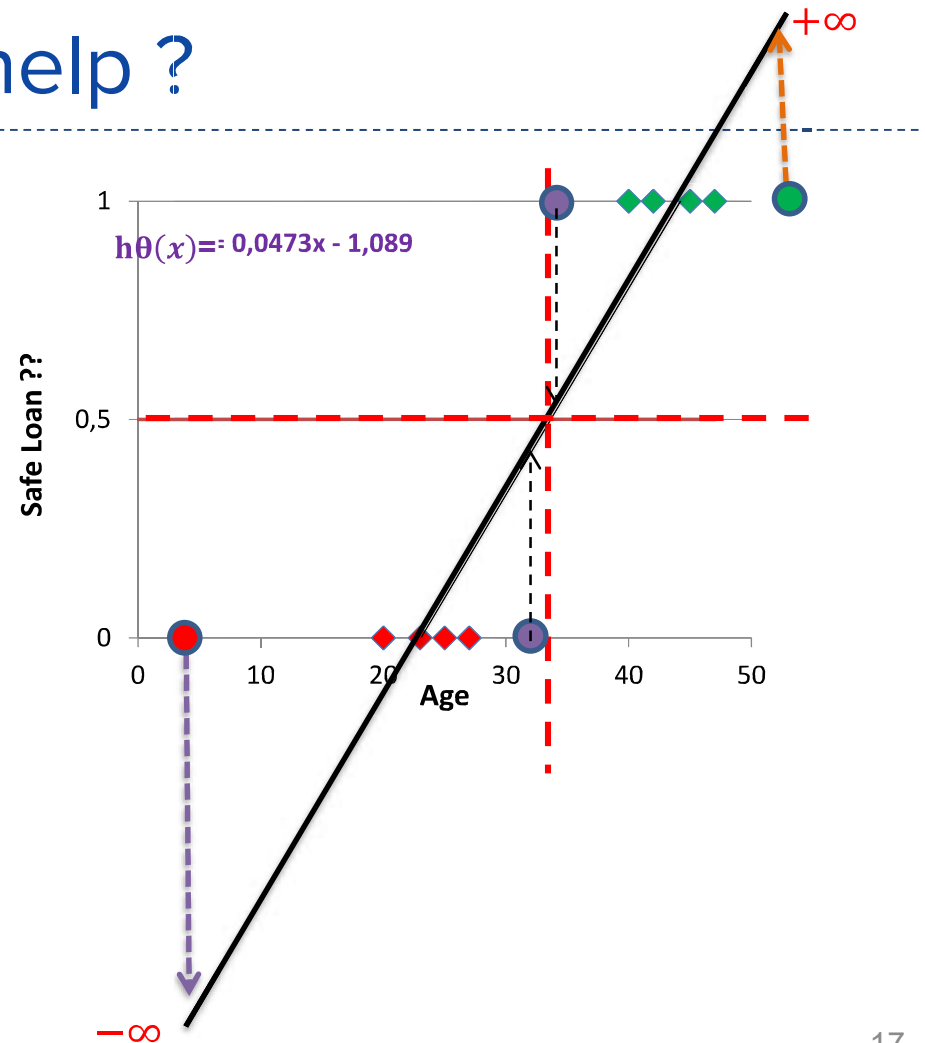
- $y = 0$ or 1
- But with LR: $h_{\theta}(x)$ can be > 1 or < 0

2nd



Can Linear Regression help ?

- How the model behaves with extreme points?
 - Certain predictions
 - And with middle points ?
 - Not very certain
 - We need to know how confident our prediction is
- 3rd
- Need for another Linear Classifier !!
→ **Logistic Regression** $0 \leq h_{\theta}(x) \leq 1$



3.3

Logistic Regression



3.3.1

Logistic Regression

intuition

Logistic Regression Model


- Change the form for our hypotheses $h_{\theta}(x) = \theta^T x$ to satisfy $0 \leq h_{\theta}(x) \leq 1$



- Use the Sigmoid / Logistic Function : $g(z) = \frac{1}{1 + e^{-z}}$
- $h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$

Interpretation of $h_{\theta}(x) = g(\theta^T x)$

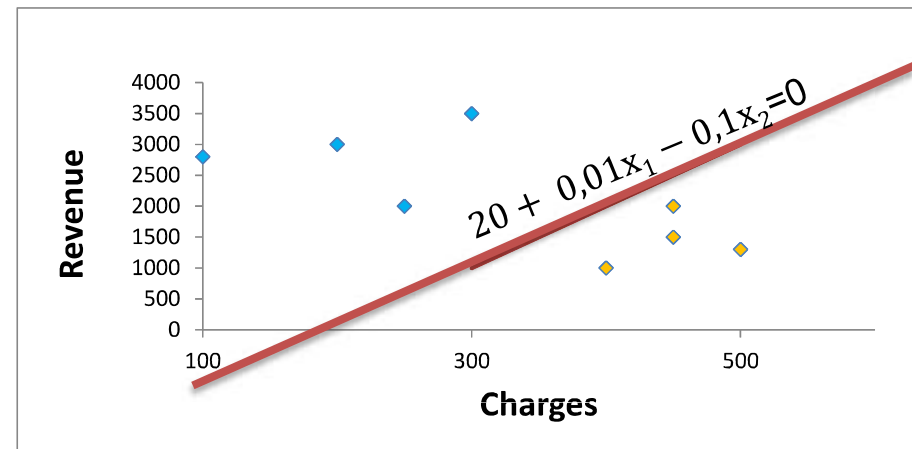
- $h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$
- $h_{\theta}(x) = \text{estimated probability that } y = 1$ given the input x parameterized by θ
 - Example: $h_{\theta}(x) = 0,8 \rightarrow$ The probability that the loan is safe ($y=1$) is equal to 80%
- $h_{\theta}(x) = P(y = 1 | x; \theta) = 1 - P(y = 0 | x; \theta)$
- $P(y = 1 | x; \theta) + P(y = 0 | x; \theta) = 1$
- Hypothesis:
 - $y = 1$ when $g(\theta^T x) \geq 0,5 \rightarrow \theta^T x \geq 0$
 - $y = 0$ when $g(\theta^T x) < 0,5 \rightarrow \theta^T x < 0$



Decision Boundary

- Example: Loan demand evaluation model
 - Predict the loan safety class given the revenue and the charges values
 - $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2) = g(20 + 0,01 \text{ \#revenue} - 0,1 \text{ \#charges})$

- Predict **1**: if $g(\theta^T x) \geq 0,5 \rightarrow \theta^T x \geq 0$
 - if $g(20 + 0,01x_1 - 0,1x_2) \geq 0.5$
 - $\Rightarrow 20 + 0,01x_1 - 0,1x_2 \geq 0$
- Predict **0**: if $g(\theta^T x) < 0,5 \rightarrow \theta^T x < 0$
 - if $g(20 + 0,01x_1 - 0,1x_2) < 0.5$
 - $\Rightarrow 20 + 0,01x_1 - 0,1x_2 < 0$



- $20 + 0,01x_1 - 0,1x_2 = 0$ is our decision boundary


Decision Boundary

- Example: Loan demand evaluation model
- Predict the loan safety class given the revenue and the charges values
 - $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2) = g(20 + 0,01 \text{ \#revenue} - 0,1 \text{ \#charges})$

Charges	Revenue	$\theta^T x$	$g(\theta^T x)$	Safe loan ? (prediction)
500	1300	-17	4,14E-08	0
450	1500	-10	4,54E-05	0
400	1000	-10	4,54E-05	0
450	2000	-5	6,69E-03	0
250	2000	15	1,00E+00	1
200	3000	30	1,00E+00	1
300	3500	25	1,00E+00	1
100	2800	38	1,00E+00	1
450	2700	2	0,880797	1

Non-linear decision boundaries

- Example: $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$
 - $\theta = [-1, 0, 0, 1, 1]$
 - Predict 1 if $-1 + x_1^2 + x_2^2 \geq 0 \Rightarrow x_1^2 + x_2^2 \geq 1$
 - Predict 0 if $-1 + x_1^2 + x_2^2 < 0 \Rightarrow x_1^2 + x_2^2 < 1$
- As with polynomial regression, we can have more complex decision boundaries by adding higher polynomial terms
 - $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1^3 + \theta_6 x_2^3)$



3.3.2

Cost Function (1)

Cost Function 1: Intuition

- Quality Metric

Real Data

Charges	Revenue	$\theta^T x$	Safe loan ?
100	2800	38	1

Charges	Revenue	$\theta^T x$	Safe loan ?
450	1500	-10	0

Predictions ?

A good model must predict: 1

A good model must predict: 0

Pick θ to maximize:
 $P(y = 1 | x; \theta)$

That is
 $P(y = 1 | x_1 = 100, x_2 = 2800; \theta)$

Pick θ to maximize:
 $P(y = 0 | x; \theta)$

That is
 $P(y = 0 | x_1 = 450, x_2 = 1500; \theta)$

Cost Function 1: Intuition

- Cost function for Linear regression


$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Charges	Revenue	$\theta^T x$	y = Safe loan (real)	predicion
500	1300	-17	0	0
450	1500	-10	0	0
400	1000	-10	0	0
450	2000	-5	0	1
250	2000	15	1	1
200	3000	30	1	1
300	3500	25	1	0
100	2800	38	1	1


Cost (error)
$h_{\theta}(x^{(1)}) - y^{(1)}$
$h_{\theta}(x^{(i)}) - y^{(i)}$
...
...

$$\frac{1}{1 + e^{-\theta^T x}}$$

Not linear



We need a cost fuction with a convex shape



Cost Function 1: Intuition

- Cost function for Logistic regression model

Charges	Revenue	$\theta^T x$	y =Safe loan (real)	predicion	Cost (error)
500	1300	-17	0	0	$\text{Cost}(h_{\theta}(x^{(1)}), y^{(1)})$
450	1500	-10	0	0	...
400	1000	-10	0	0	...
450	2000	-5	0	1	...
250	2000	15	1	1	$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$
200	3000	30	1	1	...
300	3500	25	1	0	...
100	2800	38	1	1	...

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Cost= 0

Impose high penalties (costs)

Cost Function 1: Intuition (for one data point)

$\text{Cost}(h_{\theta}(x), y) = \underbrace{-(y) \cdot \log(h_{\theta}(x))}_{\text{Predicted value}} \underbrace{-(1 - y) \cdot \log(1 - h_{\theta}(x))}_{\text{Real value}}$

If $y = 0$: $\text{Cost}(h_{\theta}(x), y) = -(1 - y) \cdot \log(1 - h_{\theta}(x))$





If $y = 1$: $\text{Cost}(h_{\theta}(x), y) = -(y) \cdot \log(h_{\theta}(x))$

$h_{\theta}(x) = 0 \Rightarrow \text{Cost}(h_{\theta}(x), y) = 0$

$h_{\theta}(x) = 1 \Rightarrow \text{Cost}(h_{\theta}(x), y) = \infty$ **High penalty**

$h_{\theta}(x) = 1 \Rightarrow \text{Cost}(h_{\theta}(x), y) = 0$

$h_{\theta}(x) = 0 \Rightarrow \text{Cost}(h_{\theta}(x), y) = \infty$ **High penalty**



Cost Function 1: Intuition (for all data points)

- For one data point: $\text{Cost}(h_{\theta}(x), y) = -(y) * \log(h_{\theta}(x)) - (1 - y) * \log(1 - h_{\theta}(x))$
- For all data points: $J(\theta) = \sum_{i=1}^m \text{Cost}(h_{\theta}(x), y)$
$$= \frac{1}{m} \sum_{i=1}^m (-(y^{(i)}) * \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) * \log(1 - h_{\theta}(x^{(i)})))$$
- To find the optimal values of θ : $\min_{\theta} J(\theta)$
- Gradient descent:
 - Initiate with a random set of values for θ
 - At each iteration: update the values of θ
 - For each feature x_j , $\theta_j = \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$

3.3.2

Cost Function (2)

Cost Function 2: Maximum-Likelihood

- Quality Metric

Real Data

Charges	Revenue	$\theta^T x$	Safe loan ?
100	2800	38	1

Charges	Revenue	$\theta^T x$	Safe loan ?
450	1500	-10	0

Predictions ?

A good model must predict: 1

A good model must predict: 0

Pick θ to maximize:
 $P(y = 1 | x; \theta)$

That is
 $P(y = 1 | x_1 = 100, x_2 = 2800; \theta)$

Pick θ to maximize:
 $P(y = 0 | x; \theta)$

That is
 $P(y = 0 | x_1 = 450, x_2 = 1500; \theta)$

Cost Function 2: Maximum-Likelihood

- Quality Metric: Maximizing likelihood, i.e the probability of good predictions

Charges	Revenue	$\theta^T x$	y =Safe loan	Choose θ to maximize
500	1300	-17	0	$P(y = 0 x; \theta)$
450	1500	-10	0	$P(y = 0 x; \theta)$
400	1000	-10	0	$P(y = 0 x; \theta)$
450	2000	-5	0	$P(y = 0 x; \theta)$
250	2000	15	1	$P(y = 1 x; \theta)$
200	3000	30	1	$P(y = 1 x; \theta)$
300	3500	25	1	$P(y = 1 x; \theta)$
100	2800	38	1	$P(y = 1 x; \theta)$

- Maximize function over all possible $\theta_0, \theta_1, \theta_2$:

$$\max_{\theta_0, \theta_1, \theta_2} \prod_{i=0}^m P(y_i | x; \theta)$$

3.4

Multiclass classification

Multi-class Classification: Example

- Image labelling



Input: x
Image pixels




Output: y
Object in image

Multi-class Classification: Formulation

- C possible classes: y can be 1, 2,..., C
- m data points

Data point	x_1	x_2	y
$x^{(1)}, y^{(1)}$	1	4	▲
$x^{(2)}, y^{(2)}$	3	1	♥
$x^{(3)}, y^{(3)}$	3	3	◆
$x^{(4)}, y^{(4)}$	4	4	◆
.....			



Learn:


$$P(y = \text{▲} \mid x; \theta)$$

$$P(y = \text{♥} \mid x; \theta)$$

$$P(y = \text{◆} \mid x; \theta)$$

Multi-class Classification: One-vs-all (one-vs-rest)

- Transform the original classification model to C 2-class models




- On a new input, to make a prediction, pick the class that maximizes: $\max_i h_{\theta}^i(x)$

3.5

Evaluating classifiers

Evaluating classifiers: Accuracy

- $\text{Accuracy} = \frac{\text{number of data points } \textit{classified correctly}}{\text{all data points}}$
- is 99% accuracy good?
 - Can be excellent, good, mediocre, poor, terrible
 - It depends on the proportion of the classes in your dataset.



In this example, 99% of the data is labeled as the negative class.

Just by predicting everything to $y = 0$, we can have an accuracy of 99%

- Accuracy is not ideal for skewed (imbalanced) classes !!

Evaluating classifiers: Confusion matrix

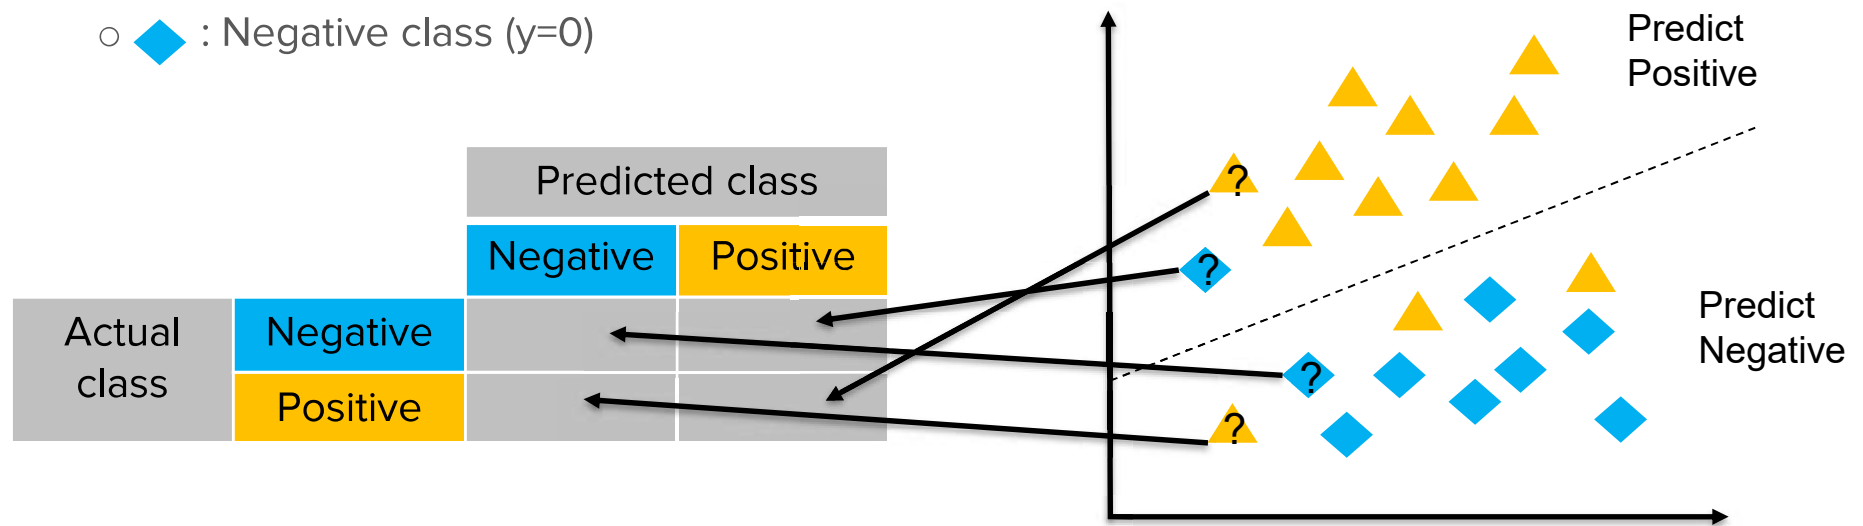
- In many cases in real life problems, you care more about well predicting one class than the others:
 - Cancer detection: care more about cancer gets detected. You can tolerate occasionally false detections but not overlooking real cancers.

Y = 0 (no cancer)	You can tolerate having errors when predicting this class: predict patient has cancer when it's not the case
Y = 1 (cancer)	You can not tolerate having errors when predicting this class: predict patient has no cancer when it's the case

- There is a need for a performance metric that can favor one type of an error than an other.

Evaluating classifiers: Confusion matrix

- We have two classes:
 - ▲ : positive class ($y = 1$)
 - ◆ : Negative class ($y = 0$)




- Match each data point to the appropriate cell

Evaluating classifiers: Confusion matrix


- Fill the cells with the corresponding numbers
 - ▲ : positive class ($y = 1$)
 - ◆ : Negative class ($y=0$)

		Predicted class	
		Negative	Positive
Actual class	Negative	?	?
	Positive	?	?



Evaluating classifiers: Confusion matrix

		Predicted class	
		Negative	Positive
Actual class	Negative	True Negative	False Positive
	Positive	False Negative	True Positive



- Sklearn: http://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html

Evaluating classifiers: Precision - Recall

		Predicted class	
		Negative	Positive
Actual class	Negative	True Negative	False Positive
	Positive	False Negative	True Positive

- Precision for the positive class answers the following question:
- Out of all the examples the classifier labeled as positive, what fraction were correct?

- $$\text{Precision} = \frac{\text{True positive}}{\text{True positive} + \text{False positive}} = \frac{9}{9+1} = 90\%$$

Evaluating classifiers: Precision - Recall

		Predicted class	
		Negative	Positive
Actual class	Negative	True Negative	False Positive
	Positive	False Negative	True Positive

- Recall for the positive class answers the following question :
- Out of all the positive examples there were, what fraction did the classifier pick up?

- $$\text{Recall} = \frac{\text{True positive}}{\text{True positive} + \text{False negative}} = \frac{9}{9+3} = 75\%$$

Evaluating classifiers: F1 score

- Given the nature of the problem, you can optimize the model to get a better precision or a better recall.
- It is possible to optimize both by combining precision and recall into a single value, called **F1 score**.
- $$F1\ score = 2 * \frac{precision * recall}{precision + recall} = 81.81\%$$
- Best value at 1, worse at 0.

3.6

K-fold cross-validation

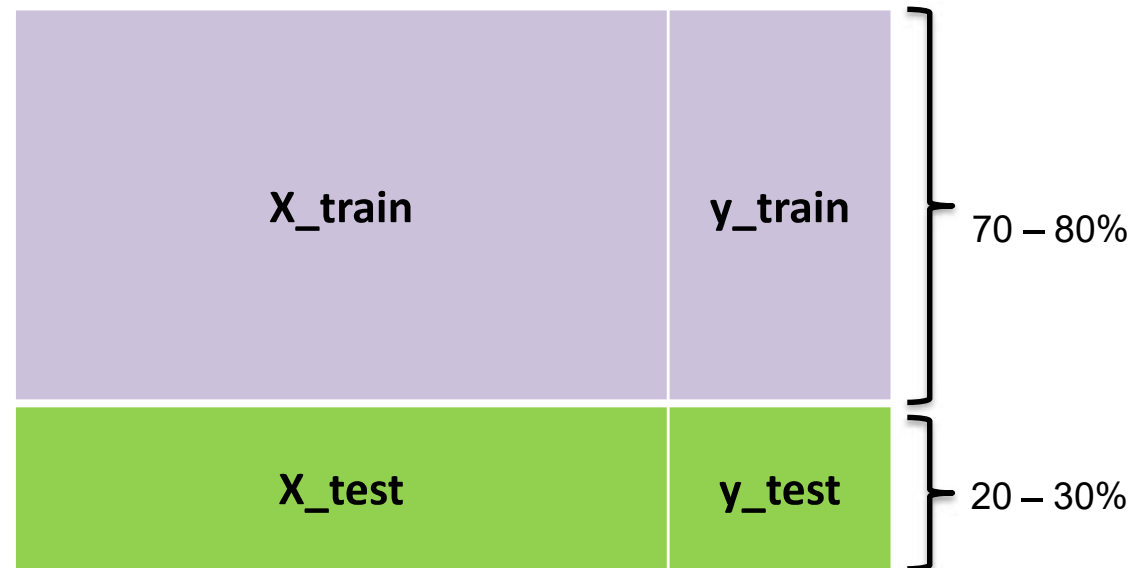
Cross-Validation: Why ?

- Data = Training set + Testing set

One-round Cross Validation

We need to derive the most accurate estimate of model prediction performance !!!

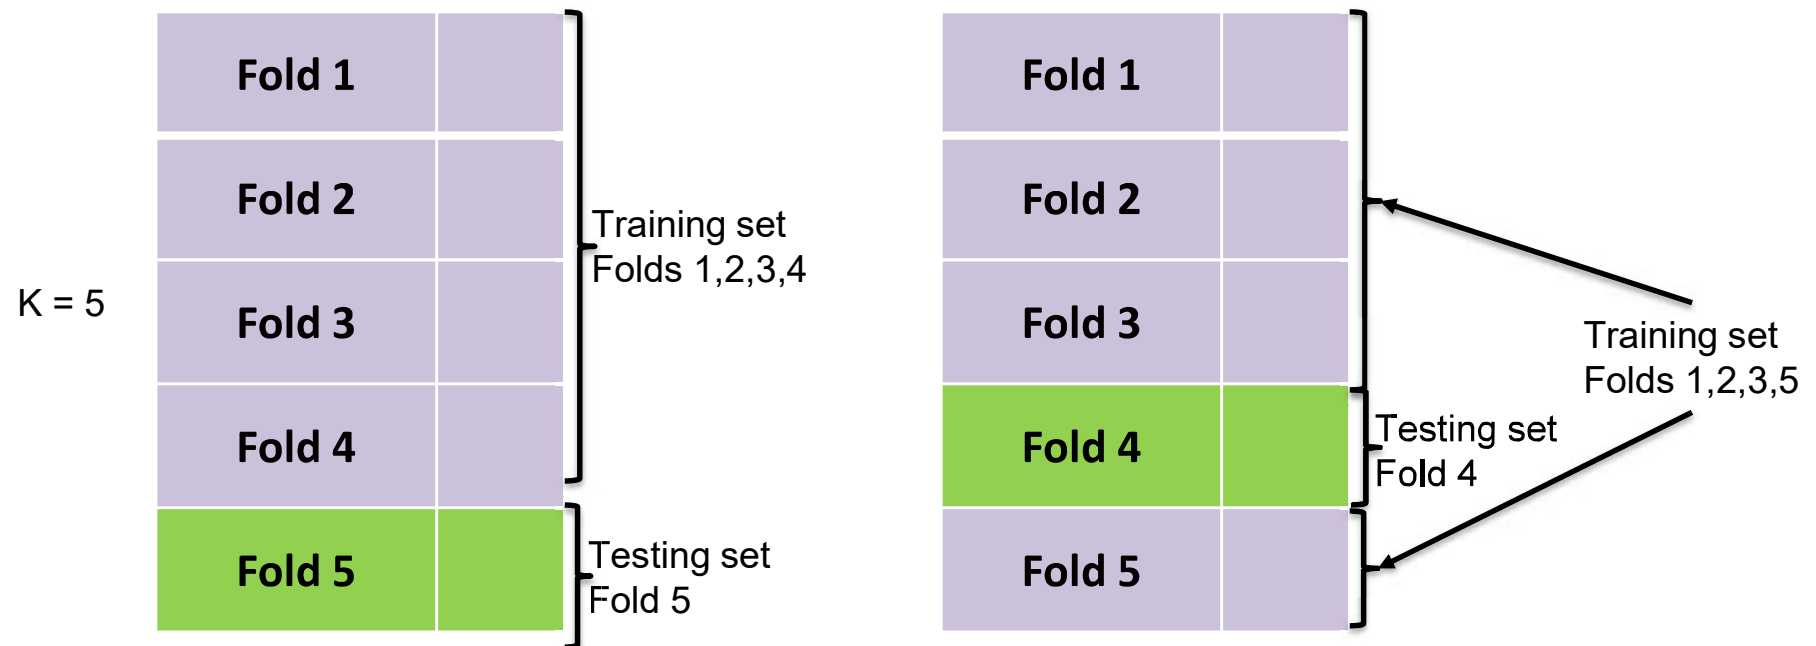
Gives an insight on how the model will generalize to an independent dataset



- Each data point is used only once for training or for testing
- Variability may arise !!
- Moreover: Sometimes we do not have enough data available to make reliable partitions

Cross Validation example: K-fold

- Partition the dataset into **K** folds (bins) of equal size.



- for each $k = 1, 2, \dots, K$, fit the model to the other $K - 1$ parts and compute its error in predicting the k^{th} part.

Cross Validation example: K-fold

- Run **K** separate learning experiments.
 - Pick testing set
 - Train
 - Test on testing test and compute performance
 - Example: Linear Regression : R^2 · Logistic Regression: Accuracy
- Average the performance from those **K** experiments
- Typically, K=5 or 10
- K-Fold is more robust for parameter tuning (choose the regularization parameter, the learning rate, ..)

More about Cross-Validation

- Multiple rounds of cross-validation are performed using different partitions, and the validation results are averaged over the rounds
- Common Types of Cross-Validation:
 - Non-exhaustive cross-validation
 - k -fold cross-validation
 - 2-fold cross-validation
 - Exhaustive cross-validation
 - Leave-p-out cross-validation
 - Leave-one-out cross-validation



**Thank you for your
attention**